

COMP 543: Tools & Models for Data Science

Relational Algebra

Chris Jermaine & Risa Myers

Rice University



Relational Calculus vs. Algebra

- In Relational Calculus
 - You say what you want
 - And not how to compute it
- But obviously...
 - This needs to be compiled into an actual computational plan
 - And in relational DBs, the plan is expressed in relational algebra
- RA is the “abstract machine” of relational databases

What Is An Algebra?

- Many Definitions!
 - Simplest: it is a set (domain) with a number of operations
 - The domain is closed under those operations
- In RA...
 - The domain is the set of all valid relations
 - The set of operations includes $\pi, \sigma, \times, \bowtie, \cup, \cap, -$
- Now let's go through the operations!

Projection

- Projection removes attributes
- $\pi_A(R)$...
 - A is a set of attributes of relation R
 - This simply removes all attributes not in A from R
 - Note: cardinality of output can differ from R
 - Output is a relation

FREQUENTS(DRINKER, CAFE)

FREQUENTS

DRINKER	CAFE
Risa	JL
Risa	BH
Chris	BH
Chris	DT

$\pi_{\text{DRINKER}}(\text{FREQUENTS})$

DRINKER
Risa
Chris

Selection

- Selection removes tuples
- $\sigma_B(R)$...
 - B is a boolean predicate that can be applied to a single tuple from R
 - This simply removes all tuples not accepted by B
 - Again: output is a relation

FREQUENTS

DRINKER	CAFE
Risa	JL
Risa	BH
Chris	BH
Chris	DT

$\sigma_{\text{DRINKER}='Risa'}(\text{FREQUENTS})$

DRINKER	CAFE
Risa	JL
Risa	BH

Selection/Projection Example

LIKES (DRINKER, COFFEE)

FREQUENTS (DRINKER, CAFE)

SERVES (CAFE, COFFEE)

? Query: Who likes 'Cold Brew' coffee?

Selection/Projection Example

LIKES (DRINKER, COFFEE)

FREQUENTS (DRINKER, CAFE)

SERVES (CAFE, COFFEE)

- Query: Who likes 'Cold Brew' coffee?

- $\pi_{\text{DRINKER}}(\sigma_{\text{COFFEE}=\text{'Cold Brew'}}(\text{LIKES}))$

Join: Cartesian Product

- Join combines tuples
- Simplest join is Cartesian product (aka: cross product)
- $R \times S \dots$
 - Returns $r \bullet s$ for all $r \in R, s \in S$
 - ? What is the output cardinality?

LIKES (DRINKER, COFFEE)
FREQUENTS (DRINKER, CAFE)
SERVES (CAFE, COFFEE)

- Often you want $\sigma_B(R \times S)$
- Shorthand for this is $R \bowtie_B S$
- ? Query: Who likes a coffee that 'Risa' likes?

Join: Theta Join

LIKES (DRINKER, COFFEE)
FREQUENTS (DRINKER, CAFE)
SERVES (CAFE, COFFEE)

- Often you want $\sigma_B(R \times S)$
- Shorthand for this is $R \bowtie_B S$
- Query: Who likes a coffee that 'Risa' likes?
 - $\text{TEMP}(d_1, c_1, d_2, c_2)$
 $\leftarrow \text{LIKES} \bowtie_{\text{COFFEE}=\text{COFFEE}} (\sigma_{\text{DRINKER}=\text{'Risa'}}(\text{LIKES}))$
 - $\pi_{d_1}(\text{TEMP})$

Join: Natural Join

LIKES (DRINKER, COFFEE)

FREQUENTS (DRINKER, CAFE)

SERVES (CAFE, COFFEE)

- Often you want to join two relations
 - Using an equality check on all attributes having the same name
 - Then project away redundant attributes
- Shorthand for this is $R * S$
- ? Query: Who goes to a cafe serving a coffee that they like?

Join: Natural Join

LIKES (DRINKER, COFFEE)

FREQUENTS (DRINKER, CAFE)

SERVES (CAFE, COFFEE)

- Often you want to join two relations
 - Using an equality check on all attributes having the same name
 - Then project away redundant attributes
- Shorthand for this is $R * S$
- Query: Who goes to a cafe serving a coffee that they like?
 - $\pi_{\text{DRINKER}}(\text{LIKES} * \text{FREQUENTS} * \text{SERVES})$

Set-Based Operations

LIKES (DRINKER, COFFEE)

FREQUENTS (DRINKER, CAFE)

SERVES (CAFE, COFFEE)

- Can use standard set operations as well: $\cup, \cap, -$
 - To use, types and numbers of input attributes must match
 - By convention, attribute names come from LHS
 - $R \cup S$: all tuples in R or in S
 - $R \cap S$: all tuples in R and in S
 - $R - S$: all tuples in R and not in S

? Query: Who does not like 'Cold Brew' coffee?

Set-Based Operations

LIKES (DRINKER, COFFEE)

FREQUENTS (DRINKER, CAFE)

SERVES (CAFE, COFFEE)

- Can use standard set operations as well: $\cup, \cap, -$
 - To use, types and numbers of input attributes must match
 - By convention, attribute names come from LHS
 - $R \cup S$: all tuples in R or in S
 - $R \cap S$: all tuples in R and in S
 - $R - S$: all tuples in R and not in S
- Query: Who does not like 'Cold Brew' coffee?
 - $\text{COLDBREWGOOD} \leftarrow \pi_{\text{DRINKER}}(\sigma_{\text{COFFEE}=\text{'Cold Brew'}}(\text{LIKES}))$
 - $(\pi_{\text{DRINKER}}(\text{FREQUENTS})) - \text{COLDBREWGOOD}$

Complicated Set-Based Example

- Query: Who does not like 'Cold Brew' coffee?

- $\text{COLDBREWGOOD} \leftarrow \pi_{\text{DRINKER}}(\sigma_{\text{COFFEE}=\text{'Cold Brew'}}(\text{LIKES}))$

- $(\pi_{\text{DRINKER}}(\text{FREQUENTS})) - \text{COLDBREWGOOD}$

? Why use FREQUENTS instead of LIKES?

Complicated Set-Based Example

LIKES (DRINKER, COFFEE)

FREQUENTS (DRINKER, CAFE)

SERVES (CAFE, COFFEE)

? Who only goes to cafes where they can get a coffee they like?

Complicated Set-Based Example

LIKES (DRINKER, COFFEE)

FREQUENTS (DRINKER, CAFE)

SERVES (CAFE, COFFEE)

- Who only goes to cafes where they can get a coffee they like?
 - Use 'all people' – 'those who go to a cafe where they can't get a coffee they like'
 - $ALLPEEPS \leftarrow \pi_{DRINKER}(FREQUENTS)$
 - How about 'those who go to a cafe where they can't get a coffee they like'?
 - Use FREQUENTS – DRINKER, CAFE 'combos where the person can get a coffee they like'
 - $GOODCOFFEE \leftarrow \pi_{DRINKER,CAFE}(LIKES * SERVES)$
- Then the answer is
 - $ALLPEEPS - \pi_{DRINKER}(FREQUENTS - GOODCOFFEE)$

Questions?