

COMP 543: Tools & Models for Data Science

Big Data Part One: An Intro to MapReduce

Chris Jermaine & Risa Myers

Rice University

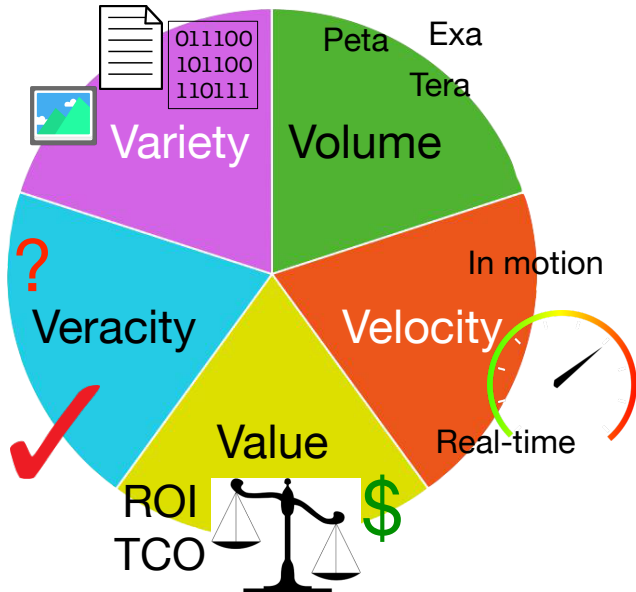


15 Years Ago...

- Say you had a big data set, wanted platform to analyze it
- What is “big”?
 - Too large to fit in RAM of an expensive server machine
 - “Big Data” termed appeared in late 90’s
- Is 5GB in 2002, a couple of terabytes in 2018
- You might spend \$20 - 30K on a server

The 5 V's of "Big" Data

- 1 Volume
- 2 Variety
- 3 Velocity
- 4 Veracity
- 5 Value



7 V's

- 1 Volume
- 2 Variety
- 3 Velocity
- 4 Veracity
- 5 Value
- 6 Variability
- 7 Visualization

How to Analyze a “Big” Dataset

- You might write your own software
- Costly, time consuming
 - A \$10M software feature might eat up most of the IT budget for a single firm
- Requires expertise not always found in house
- Risky: high potential for failure

Or, You Might Buy a DB System

- Costs a LOT of money
- Performance often unpredictable, or just flat out poor
- Software can be insanely complicated to use correctly
- Software stack too big/deep, not possible to unbundle
 - If you are doing analysis, ACID not important
 - And yet, you pay for it (money, complexity, performance)
- Difficult to put un- or semi-structured data into an SQL DB

Plus, Many People Just Don't Like SQL

- People uncomfortable with declarative programming
 - We love it!
 - But users don't really know what's happening under the hood
 - Makes many programmers uncomfortable
- Also, not easy/natural to specify important computations
 - Especially data mining and machine learning
 - Not to mention High Performance Computing-style computations (Analytics)

By Early-Mid 2000's...

- The Internet companies (Google, Yahoo, etc.)...
 - ...had some of the largest databases in the world
 - But they never used classical SQL databases for webscale
- How'd they do it?
 - Many ways...
 - But paradigm with most widespread impact was MapReduce
 - First described in a 2004 academic paper, appeared in Operating Systems Design and Implementation (OSDI)
 - “MapReduce: Simplified Data Processing on Large Clusters”

Stage is set for a New Paradigm

- Existing approaches don't work
 - Custom software is too expensive
 - RDBMSs are too rigid
- Data volume has exploded
- Commodity hardware has become cheap
- Cost is proportional to workload
- Welcome MapReduce!

What Is MapReduce?

- A simple data processing paradigm
- Leverages a cluster of commodity machines
 - For distributed processing
 - For fault tolerance
 - Requires a distributed file system
 - See “The Google File System” by Ghemawat, et al. 2003
- Well suited to calculations that can be computed in independent chunks

What Is Novel about MapReduce?

- Novelty is in scalability and fault tolerance
- The functional model already existed
- Consider the map and reduce functions in Python

- Map:

```
items = [ 1, 2, 3, 4, 5 ]  
mod2 = list(map(lambda x:  x % 2, items))  
[1, 0, 1, 0 , 1]
```

- Reduce:

```
from functools import reduce  
mySum = reduce((lambda x,y:  x + y), items)  
15
```

- To process a data set:
 - You have two pieces of user-supplied code
 - A Map code
 - And a Reduce code

- These are run in a huge shared-nothing compute cluster
- Using three data processing phases
 - 1 A Map phase
 - 2 A Shuffle phase
 - 3 And a Reduce phase

First: What Is Shared-Nothing?

- Store/analyze data on a large number of commodity machines
 - Local, non-shared storage attached to each of them
 - Only link is via a LAN
 - Shared nothing refers to no sharing of RAM, storage
 - Note: Network Attached Storage (NAS) is common now, “pure” shared-nothing rarer

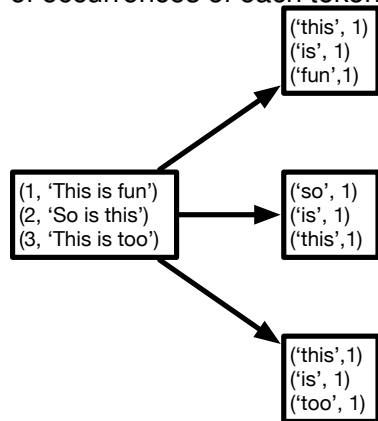
Benefits of Shared-Nothing

- Inexpensive, built out of commodity components
- Compute resources scales nearly linearly with money
- Contrast to shared RAM machine with uniform memory access
- Easier to program than Shared Memory systems

MapReduce: The Map Phase

- Input data are stored in a huge file
 - Contains a simple list of pairs of type $(key1, value1)$
- Have a UDF of the form $Map(key1, value1)$
 - Outputs a list of pairs of the form $(key2, value2)$
- During the Map phase of the MapReduce computation
 - The *Map* function is called for every record in the input
 - Instances of *Map* run in parallel all over the cluster

Task: Compute the number of occurrences of each token

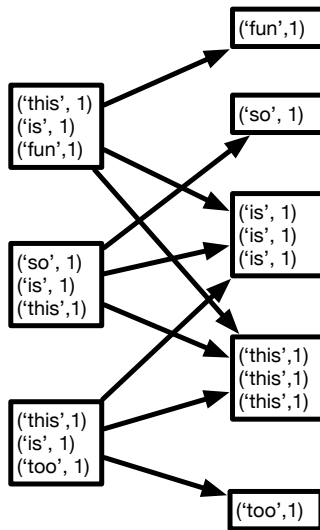


Example: Word Count

- Large text corpus
- Want to count number of occurrences of each word
- Ex output: ('The', 1832321), ('An', 1732432), etc.
- To power the Map phase, MapReduce software automatically:
 - Breaks the corpus into large number of $(lineNo, text)$ pairs
 - Distributes the Map UDF
 - Distributes the fragments
 - Ensures that the UDF is run on all the fragments

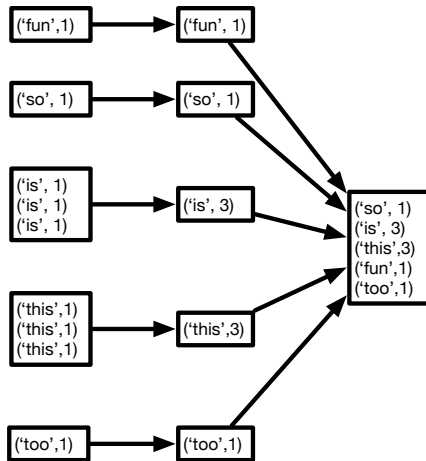
MapReduce: The Shuffle Phase

- Accepts all of the $(key2, value2)$ pairs from the Map phase
 - And it groups them together
- After grouping, all of the pairs
 - From all over the cluster having the same $key2$ value
 - Are merged into a single $(key2, itemize\langle value2 \rangle)$ pair
- Called a “Shuffle”...
 - Because this is where a potential all-to-all data
 - Can be expensive
 - Often implemented using a hash function to map keys to nodes transfer happens

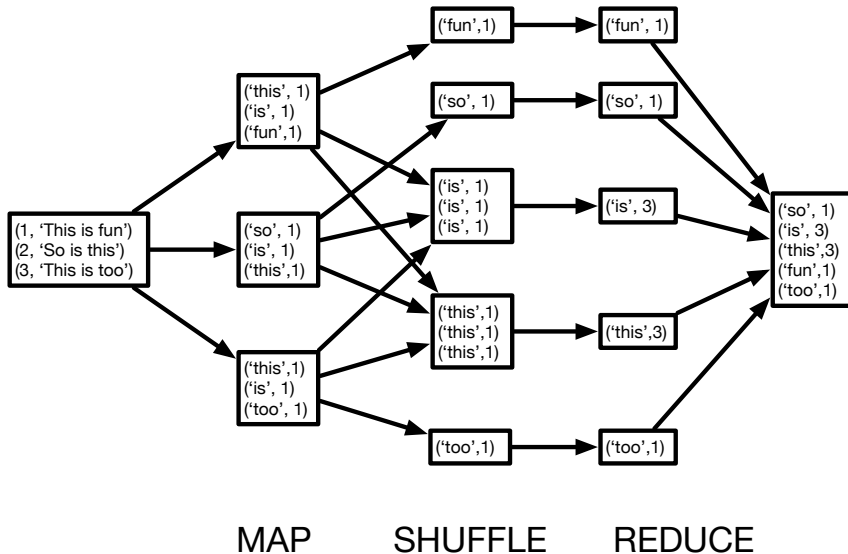


MapReduce: The Reduce Phase

- Have a user-supplied function of the form
 - $Reduce(key2, itemize\langle value2 \rangle)$
 - Outputs a list of $value3$ objects
- In the Reduce phase of the MapReduce computation
 - $Reduce$ function is called for every $key2$ value output by the Shuffle
 - Instances of $Reduce$ run in parallel all over the compute cluster
 - The output of all of those instances is collected
 - Put in a (potentially) huge output file



MapReduce: All Phases



MapReduce is a Compute Paradigm

- It is not a data storage paradigm
 - But must read/write data from some storage system
- So MapReduce is strongly linked with the idea of a distributed file system (DFS)
 - Allows data to be stored/accessed across machines in a network
 - Abstracts away differences between local and remote data
 - Uses the same API to read/write data
 - Regardless of network location

- DFSs have been around for a long time
 - First widely used DFS was Sun's NFS, first introduced in 1985
 - Intended to be allow file access across networks as if file was local
- Unlike classical DFS...
 - MapReduce DFS sits on top of each machine's OS
 - Lives in "user space"
 - The OS is not aware of the DFS
 - You can't mount it anywhere
- Why on top of, not in the OS?
 - Heterogeneity no problem (each machine can run a different OS)
 - Easily portable (JVM)
- Even as MapReduce becomes less popular, MR DFS lives on!

- MapReduce pros
 - MUCH lower programmer burden than HPC
 - No synchronization, all parallelism is implicit
 - Data and task placement automatic
 - Built-in fault tolerance
 - Works with (almost!) arbitrarily-sized data
 - Out-of-core execution is no problem

- MapReduce cons
 - Standard softwares are JVM-based
 - Not suitable for communication-heavy tasks...
 - ...only communication is via the shuffle
 - Assumes BIG data... always reads/writes data from DFS

Questions?