

COMP 543: Tools & Models for Data Science

Big Data: An Intro to MapReduce

Chris Jermaine & Risa Myers

Rice University



15 Years Ago...

- Say you had a big data set, wanted platform to analyze it
- What is “big”?
 - Too large to fit in RAM of an expensive server machine
- Is 5GB in 2002, a couple of terabytes in 2018

You Would Buy a DB System

- Cost A LOT of money
- Performance often unpredictable, or just flat out poor
- Software insanely complicated to use correctly
- Software stack too big/deep, not possible to unbundle
 - If you are doing analysis, ACID not important
 - And yet, you pay for it (money, complexity, performance)
- Difficult to put un- or semi-structured data into an SQL DB

Or You Roll Your Own Software

- Costly, time consuming
 - A \$10M software feature might eat up most of the IT budget for a single firm
- Requires expertise not always found in house
- Risky: high potential for failure

Many People Just Don't Like SQL

- People uncomfortable with declarative programming
 - We love it!
 - But user doesn't really know what's happening under the hood
 - Makes many programmers uncomfortable
- Also, not easy/natural to specify important computations
 - Especially data mining and machine learning

By Early-Mid 2000's...

- The Internet companies (Google, Yahoo, etc.)...
 - ...had some of the largest databases in the world
 - But they never used classical SQL databases for webscale
- How'd they do it?
 - Many ways...
 - But paradigm with most widespread impact was MapReduce
 - First described in a 2004 academic paper, appeared in OSDI
 - Easy read! Do a search on "Google MapReduce paper"

What Is MapReduce?

- It is a simple data processing paradigm
- To process a data set:
 - You have two pieces of user-supplied code
 - A Map code
 - And a Reduce code
- These are run in a huge shared-nothing compute cluster
 - Using three data processing phases
 - A Map phase
 - A Shuffle phase
 - And a Reduce phase

First: What Is Shared-Nothing?

- Store/analyze data on a large number of commodity machines
 - Local, non-shared storage attached to each of them
 - Only link is via a LAN
 - Shared nothing refers to no sharing of RAM, storage
- Why good?
 - Inexpensive, built out of commodity components
 - Compute resources scales nearly linearly with money
 - Contrast to shared RAM machine with uniform memory access

MapReduce: The Map Phase

- Input data are stored in a huge file
 - Contains a simple list of pairs of type (key1,value1)
- Have a UDF of the form Map(key1,value1)
 - outputs a list of pairs of the form (key2, value2)
- In the Map phase of the MapReduce computation
 - The Map function is called for every record in the input
 - Instances of Map run in parallel all over the cluster

MapReduce: The Shuffle Phase

- Accepts all of the (key2, value2) pairs from the Map phase
 - And it groups them together
- After grouping, all of the pairs
 - From all over the cluster having the same key2 value
 - Are merged into a single (key2, list <value2>) pair
- Called a Shuffle...
 - Because this is where a potential all-to-all data transfer happens

MapReduce: The Reduce Phase

- Have a user-supplied function of the form
 - Reduce (key2,list <value2>)
 - Outputs a list of value2 objects
- In the Reduce phase of the MapReduce computation
 - Reduce function is called for every key2 value output by the Shuffle
 - Instances of Reduce run in parallel all over the compute cluster
 - The output of all of those instances is collected
 - Put in a (potentially) huge output file

MapReduce Is a Compute Paradigm

- It is not a data storage paradigm
 - But any MapReduce system
 - Must read/write data from some storage system
- So MapReduce strongly linked with the idea of a distributed file system (DFS)
 - Allows data to be stored/accessed across machines in a network
 - Abstracts away differences between local and remote data
 - Same API to read/write data
 - No matter where data is located in the network

Distributed File Systems for MR

- DFSs have been around for a long time
 - First widely used DFS was Sun's NFS, first introduced in 1985
- Unlike classical DFS, MapReduce DFS sits on top of each machine's OS
 - Lives in "user space"
 - The OS is not aware of the DFS
 - You can't mount it anywhere
- Why on top of, not in the OS?
 - Hetrogeneity no problem
 - Easily portable (JVM)
- Even as MapReduce becomes less popular, MR DFS lives on!

Questions?