

# Tools & Models for Data Science

## Linear Regression


Chris Jermaine & Risa Myers

Rice University



# Linear Regression

- Most common model in data science! (Logistic Regression is very common as well)
  - Have a set of training data
  - Bunch of  $(x_i, y_i)$  pairs
  - $x_i$  is a vector of real-valued “regressors” / features / dimension
  - $y_i$  is a real-valued “response”
  - Want to learn a model that, given a new  $x$ , can predict  $y$



- Model is exceedingly simple

- Predict  $\hat{y}$  as  $x \cdot r$
- $r$  is a vector of “regression coefficients”
- $x \cdot r$  is the dot product of:  $x \cdot r = \sum_j x_j \times r_j = \hat{y}_i$
- Can be used with loss functions:

- Least Squares (L2 norm)

$$Loss = \|y - f(x)\|_2^2$$

- Mean Squared Error (MSE), where  $n$  is the number of training points

$$Loss = \frac{\|y - f(x)\|_2^2}{n}$$

- Others...

# Regression Coefficient Example

- Specify the weight/importance and direction of each feature
- Weight is indicated with magnitude
- Direction is indicated by the sign

Predicting Song Tempo

<b>Feature</b>	<b>Regression Coefficient value</b>
Duration	-0.0061
Latitude	-0.1197
Loudness	1.1527
Year	0.0013
Intercept	139.72

# Our Data

- Let the matrix  $\mathbf{X}$  store the training data
- $i$ th row in  $\mathbf{X}$  is  $i$ th training point,  $x_i$
- $\mathbf{y}$  is a column vector storing responses

$$\mathbf{X} = \begin{bmatrix} \text{---} & x_1 & \text{---} \\ \text{---} & x_2 & \text{---} \\ & \vdots & \\ \text{---} & x_i & \text{---} \\ & \vdots & \\ \text{---} & x_n & \text{---} \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_i \\ \vdots \\ y_n \end{bmatrix}$$

# How to Learn?

- Turns out there is a closed-form solution to this minimization problem we are solving for regression
  - Then closed form least-squares estimate for  $r$  is (you can look this up):

$$\hat{r} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- This minimizes loss:

$$\sum_i (y_i - x_i \times r)^2$$

$$\hat{r} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- But this can be problematic for “Big Data”... why?

## Problematic for “Big Data”

- [illegible]



# More Reasonable Big Data Formulation

- Recall the closed form least squares estimator for  $\hat{r}$ :

$$\hat{r} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- We can compute  $(\mathbf{X}^T \mathbf{X})^{-1}$  as:

$$\left( \sum_i x_i^T x_i \right)^{-1}$$

- Note: assumes  $x_i$  is a row vector
- $x_i^T x_i$  is the outer product of  $x_i$  with itself, resulting in an  $n \times n$  matrix
- Recall from lab that  $\sum_i x_i^T x_i$  is the sum of the outer products of the matrix rows

? What's great about this formulation?

- Compute  $(\mathbf{X}^T \mathbf{X})^{-1}$  as:

$$\left( \sum_i x_i^T x_i \right)^{-1}$$

- What's great about this formulation?
  - It can be parallelized!
  - Distribute blocks of rows (say 100) at a time
  - Compute the products
  - Collect, reassemble, sum, then invert

# More Reasonable Big Data Formulation (continued)

- Goal: Compute the closed form least squares estimate for  $\hat{r}$

$$\hat{r} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- 1 Compute  $(\mathbf{X}^T \mathbf{X})^{-1}$  as (per the last slide):

$$\left( \sum_i x_i^T x_i \right)^{-1}$$

- 2 Compute  $\mathbf{X}^T \mathbf{y}$  as:

$$\left( \sum_i (x_i \times y_i) \right)^T$$

- #2 can also be parallelized
- since it is a sum of products

$$\left( \sum_i x_i^T x_i \right)^{-1} \left( \sum_i (x_i \times y_i) \right)^T$$

? Why?

# Problematic for Very High-D Data

- Inverting  $\mathbf{X}$  can be expensive, if the number of dimensions is high
  - $\approx 100\text{K} \times 100\text{K}$  is an upper limit for a single machine
  - Laptop maxes out at  $4\text{-}8\text{K} \times 4\text{-}8\text{K}$
  - The matrix doesn't fit in memory!

? What's the solution?

# Problematic for Very High-D Data

- Closed form LR takes too much memory for High-D data
- So, don't use it!
- Instead, use Gradient Descent on the Mean Squared Error Loss function:

$$\frac{\sum_i (y_i - x_i \times r)^2}{n}$$

- where  $r$  is the vector of regression coefficients
- and  $n$  is the number of data points

- The partial derivative of the loss function wrt  $r_j$  is:

$$\frac{\partial}{\partial r_j} \frac{\sum_i (y_i - \sum_{j'} x_{i,j'} \times r_{j'})^2}{n} = \frac{\sum_i -2(y_i - \hat{y}_i)x_{i,j}}{n}$$

- Where  $\hat{y}_i$  is the prediction for  $y_i$  given the current model
- Again, this expression can be parallelized

# Gradient Descent Algorithm

$n \leftarrow$  the number of training data points

$r^0 \leftarrow$  non-stupid guess for  $r$ ;

$iter \leftarrow 1$ ;

repeat {

**for each**  $j$

$$\Delta_j \leftarrow \sum_i -\frac{2}{n} (y_i - \hat{y}_i) x_{i,j}$$

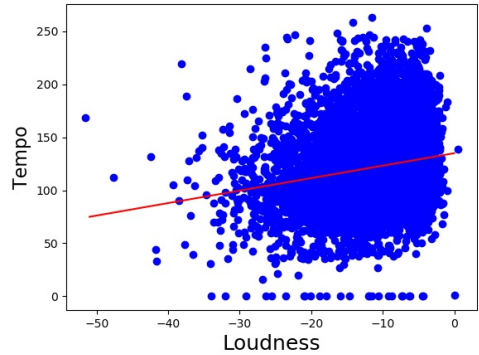
$$r^{iter+1} \leftarrow r^{iter} - \lambda \Delta;$$

$iter \leftarrow iter + 1$ ;

} while (change **in** loss  $> \epsilon$ )

- Where  $r$  is the vector of regression coefficients
- and  $n$  is the number of training data points
- Say our estimate ( $\hat{y}_i$ ) for point  $i$  is too large
- That is,  $y_i - \hat{y}_i$  is negative (example,  $-0.05$ )
- If  $x_{i,j}$  is positive (ex: 2.0), point  $i$  will try to pull  $\Delta_j$  so it is positive: contribution to  $\Delta_j$  is  $-\frac{2}{n}(-0.05)2.0 = 0.2$
- Since  $r^{iter+1} \leftarrow r^{iter} - \lambda \Delta$ , point  $i$  will try to decrease  $r_j$ : will contribute a decrease of  $\lambda 0.2$





# Why is this a Good Big Data Algorithm?

$$\frac{\partial}{\partial r_j} \frac{\sum_i (y_i - \sum_{j'} x_{i,j'} \times r_{j'})^2}{n} = \frac{\sum_i -2(y_i - \hat{y}_i)x_{i,j}}{n}$$

- It's linear in number of data points
- Also linear in number of regressors (features / dimensions)

# Why is this a Good Big Data Algorithm?

- In particular, nice for sparse data (common in really high dimensions)
  - If  $x_{i,j}$  is zero, no contribution to  $\Delta_j$
  - Note: You must use a sparse matrix representation to benefit


$$\begin{array}{l} \text{Dense:} \end{array} \begin{bmatrix} 0 & 2 & 0 & 0 \\ 0 & 0 & 7 & 0 \\ 0 & 0 & 0 & 0 \\ 4 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & -5 & 3 \end{bmatrix} \quad \begin{array}{l} \text{Sparse:} \end{array} \begin{bmatrix} 0: & (1,2) \\ 1: & (2,7) \\ 2: & \\ 3: & (0,4) \\ 4: & \\ 5: & (3,1) \\ 6: & (2,-5), (3,3) \end{bmatrix}$$

# Why is this a Bad Big Data Algorithm?

- It's linear in number of data points
- Also linear in number of regressors (features / dimensions)
- Alternatives
  - Mini-batch Gradient Descent - use a small number of randomly sampled data points at each iteration
  - Stochastic Gradient Descent - use a single randomly sampled data point at each iteration

# How To Add an Intercept?

- Add an extra column to each data point
- Always has a “1” value
- ? Why will this work?



# How To Add an Intercept?

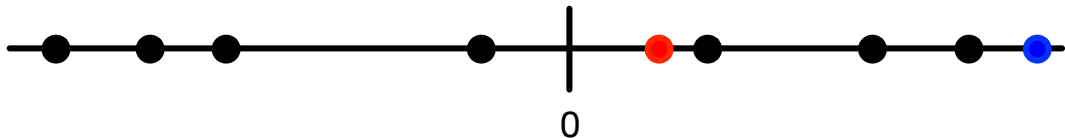
- Add an extra column to each data point
- Always has a “1” value
- Why will this work?
  - The model can learn a regression coefficient for that dimension
  - This is the intercept

# How To Handle Categorical Data?

- Easiest: during training, treat “yes” as  $+1$ , “no” as  $-1$ 
  - When applying model:  $> 0$  becomes “yes”
  - When applying model:  $< 0$  becomes “no”
- But generally this mapping is understood to leave accuracy on the table. Why?

# How To Handle Categorical Data?

- Easiest: during training, treat “yes” as  $+1$ , “no” as  $-1$
- But generally this mapping is understood to leave accuracy on the table.  
Because
  - Every “yes”/“no” treated same way
  - Tries to map all “yes” cases to  $+1$
  - Tries to map all “no” cases to  $-1$
  - Even though not all “yes” (and all “no”) cases are the same
  - The blue point is a strong “yes” than the red point

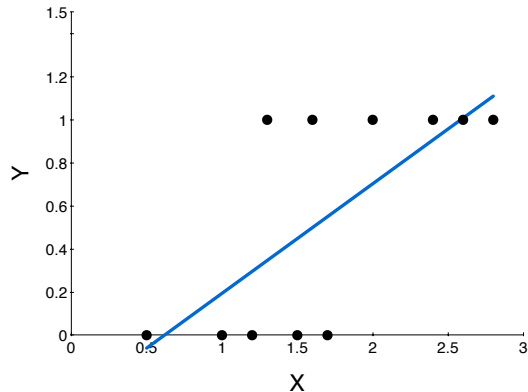
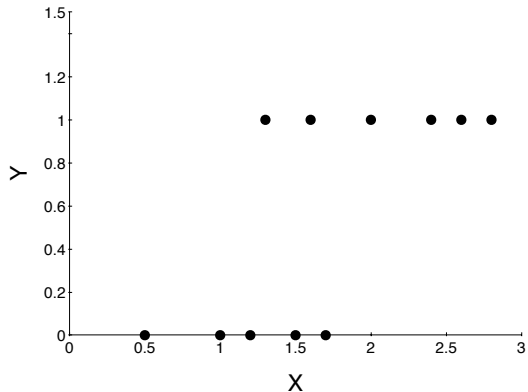






# Why is this a Problem?

- Example: A song's duration and loudness, can we predict if the tempo will be  $\geq 100$ ?
  - One song might have a really short duration, be really quiet
  - Another song might be of average duration, and be a little quiet
  - Linear Regression for **categorical** data tries map both to  $-1$ 
    - Rather than letting first map to arbitrarily large value (like  $+10$ ), a really solid "yes"
    - And letting the second map to a smaller value (like  $+0.5$ ) since a less solid "yes"
- Answer: logistic regression... will consider next time
  - Under topic of "generalized linear models"
  - Are a general class of probabilistic models based on LR
  - Logistic regression will allow more obvious "yes" cases to fall far above decision boundary
  - While obvious "no" cases fall far below

# Data not Handled Well by Linear Regression: Categorical



# Data not Handled Well by Linear Regression: Other Non-Linear



# Questions?

- What do we know now that we didn't know before?
- How can we use what we learned today?