

Tools & Models for Data Science

Neural Networks (1)

Chris Jermaine & Risa Myers

Rice University




Objectives

- Learn what an Artificial Neural Network is
- Learn the history of neural networks
- Learn components of and terminology for neural networks
- Learn the types of problems that can be addressed with neural networks


What Is a Neural Network?

- At highest level
 - It is a nonlinear function
 - Often represented as a graph



Some History: 1950s

- Interest in artificial NNs started in the 1950's
 - When simple variants were first proposed as a computational model
 - Motivated by the human brain
 - The idea was to model an animal brain
 - Where neurons are either on or off
- Interest in artificial neural networks dwindled as the von Neumann architecture won out




von Neumann architecture

Some History: 1970s

- Then back-propagation was invented in 1975
 - Effectively, it's using gradient descent to train a NN
 - Given examples, learn a model
- Took ANN further from biological roots
- But made them practically usable

Some History: 1980s

- 1980's: the computers couldn't handle the computations
- 2nd "AI winter" 1987 - 1993
 - No serious progress
 - Government felt that scientists were lying to them



Some History: 1990s


- Interest was renewed, mainly via statistical methods
 - As methods such as SVMs gained popularity
 - Kernel trick surfaced around 1992
- By late 1990's ANNs were almost seen as a joke
 - No one who was hip and with it actually looked at them
 - Why?
 - Way too data hungry
 - Way too computationally expensive to train
 - Alternatives (e.g. SVM) worked pretty well

Some History: 2000s

- But by the end of the aughts, that changed
 - We had tons of data (Google!)
 - We have tons of computational ability
 - Even though Moore's law has stalled out, we have multi-core, GPUs and the like
 - And ANNs are very amenable to parallelization
- Suddenly, ANNs were doing things that amazed everyone
 - Massive increase in accuracy in speech, image recognition
 - Learning to play Go, beat the best players in the world
 - AlphaGo Zero: Learned by playing games against itself



What Does a NN Look Like?



Terminology (1)

■ Inputs


- Data!
- # is based on the shape of your input
- Often floating point values
- Could be a multidimensional feature vector
- Could be images

■ Hidden layer(s)


- Receive input
- Perform transformation
- Pass values on
- Often fully connected

■ Outputs


- Often probabilities
- Reflect levels of certainty
- # depends on what you want e.g. a probability for each class in a set



What are the Steps in NN Processing?



Terminology (2)




Terminology (2)

- The circles are called “neurons”
 - Little computational units
 - Computed function is usually continuous and “nice”
 - Since learning is via gradient descent ¹
- Output of a neuron is called an “activation”
- Neurons are organized into “layers”
 - In simplest case, layers take input only from preceding layer
 - “Hidden” because they are not visible externally

¹There are more complex variants

Terminology (4)



- **Deep** Neural Networks have > 1 hidden layer
- Top/right layer in a NN is the “output layer” (layer t for “top”)
- Bottom/left layer in a NN is the “input layer” (layer 1)
 - Activation of input neurons read directly from input features
- One or more “hidden layers” are between those two (layers $1 < i < t$)
- Each link between neurons has a “weight”
 - Input into a neuron is the weighted sum of activations that travel over all edges that point into the neuron


Neural Network Types

- Today we'll assume a “fully connected” NN
 - That is, each neuron in layer $i > 1$ takes input from all neurons in layer $i - 1$
- But there are other connection patterns between layers ²
 - Convolutional layers
 - Used extensively in image processing
 - Where subsets of neurons are constrained to only work on contiguous regions of input data
 - Recurrent Neural Networks (RNNs)
 - In an RNN, input from upper layer fed recursively into input in lower layer
 - Input to a layer can come from > 1 layer

²later lectures


Simplest NNs

- Are called “feed-forward” NNs
 - Because activations only flow forward thru network
 - Flow from input layer
 - Thru hidden layers
 - And out of the output layer




What Function Does Each Neuron Compute?

- Need each neuron to compute some activation function $f(v)$
 - v is weighted sum over input activations
 - $f(v)$ is often simple and differentiable



Activation Functions

- Classically, the activation function is a “sigmoid” function
 - So neuron goes from “off” to “on” (like a soft binary gate)
 - Output is a floating point number
 - But smoothly, so it is differentiable (gradient descent once again!!)




Classic Sigmoid Functions

- First, the hyperbolic tangent function

$$y = \tanh(v)$$

- Second, the logistic function

$$y = (1 + e^{-v})^{-1}$$




Modern Activation Function

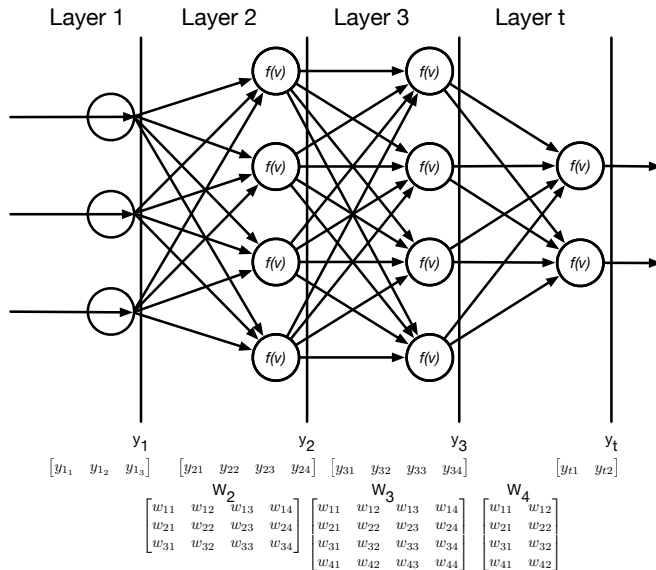
- ReLU “Rectified linear unit”

$$y = \max(0, v)$$

- 0 or pass through the input value



NN Evaluation



Evaluation of a Feed-Forward NN

- Typically done via vectorized operations... fast!
- How to push activation thru a network...
 - Let y_i be a vector denoting output of i th layer
 - Define $y_1 = x$, the input data
 - Let f denote the vector-valued function applying the activation function to each item in a vector
 - Let W_i denote the matrix of weights connecting layer $i - 1$ to layer i
 - If layer $i - 1$ has n neurons, layer i has m neurons, then W_i has n rows, m columns
 - And $y_i = f(W_i \otimes y_{i-1})$, the tensor product of W_i and y_{i-1}

$$\begin{bmatrix} & W & \end{bmatrix} \otimes \begin{bmatrix} y_{i-1} \end{bmatrix} = \begin{bmatrix} y_i \end{bmatrix}$$

Evaluation of a Feed-Forward NN

- It's all about matrix & vector multiplication
- ... and applying a non-linear activation function
- You could write lots of loops in Python
- Recall, we've talked a lot about vector operations
- GPUs are really good at linear algebra
- So, write linear algebra


In practice

- We send multiple inputs through the network at the same time
- So, y is a matrix, not a vector
- Why?
 - It's more efficient
 - Theoretically, it's the same complexity, but in practice, it's different


Bias

- Just like in Linear Regression
- Have a special neuron at each layer
- It always outputs the value 1, regardless of its inputs
- Recall i is the layer and j is the neuron index
- Then the last row of each W_i will have biases
- Since $w_{i,j} \times 1$ will be added into input to each neuron in layer $i+1$
- $y_i = f(W_i \otimes y_{i-1} + b_i)$
- b_i is a scalar that is added to every entry:


$$\begin{bmatrix} & W \\ & \end{bmatrix} \otimes y_{i-1} + b_i \begin{bmatrix} 1 \\ \cdot \\ \cdot \\ \cdot \\ 1 \end{bmatrix} = \begin{bmatrix} y_i \end{bmatrix}$$



NN Example



NN Example with Multiple Inputs




Producing Outputs

- NNs classically used for classification
 - Output is function of activations of the last hidden layer
- Often, we want certainty - e.g. a probability
- Typically we use a “softmax” function
 - That is, neuron j at layer t (the “top” layer) outputs:


$$\frac{\exp(y_{t-1,j})}{\sum_{j'} \exp(y_{t-1,j'})} = \frac{\text{output of previous layer}}{\text{normalization term}}$$

- Since all y_t values then sum to one...
 - $y_{t,j}$ typically viewed as the probability that correct class label for data y_1 is j

NN with Softmax



NN Example



Another Example Classification

- Want to solve a classification problem
- Chihuahua or muffin?
- Push the picture through the FF network
- End with two neurons
 - Chihuahua
 - Muffin



Another Example

- E-Discovery: Want to identify emails that are relevant to a court case
 - Say there are 100 topics of interest to the case
 - Crowd source relevant example emails for each topic
 - Each email might contain any number of the topics
 - Goal: identify additional emails that are relevant to each topic
- Amazon
 - Search
 - Profitability
 - Relevance
 - Third-party sellers
 - Product recommendations

Yet Another Example

- Want to identify objects in a image
- Each image may contain multiple objects
- Push the pictures through the FF network
- End with a neuron for every object you want to identify
 - Cat
 - Dog
 - Puppy
 - Tree
 - Mug
 - ...
- The output of each neuron gives the confidence of that object being in the image



- What is a NN?
- What are the components of a NN?
- What kinds of problems can it solve?

- What are the weights?
- Learned via gradient descent
- Gives rise to the famous “back-propagation” algorithm