

Tools & Models for Data Science

Deep Neural Networks (4): LSTM

Chris Jermaine & Risa Myers

Rice University




- Learn about the issues with RNNs and learn about a neural network architecture that solves some of these issues

- Have a chain of repeating “modules”
 - Each is a simple feed-forward network
 - Often a single, fully-connected layer
 - Each neuron uses some standard activation function (tanh, logistic)

Unrolling RNNs

- Training/prediction works via unrolling
- Input at each time tick concatenated with last internal state
 - Sent through a tanh/logistic layer
 - Result of that layer used to produce output, sent into next tick (Jordan network)
 - Each module is 1 iteration of the network



¹Reprinted with permission. <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Problem With RNNs

- The “vanishing gradient” problem
- During back-propagation
 - Update magnitude drops exponentially with distance from output
 - Recall

$$\frac{\frac{\partial L}{\partial w_{i,j,k}}}{\text{Neuron Output}} = \frac{\frac{\partial L}{\partial y_{i,k}}}{\text{Neuron Output}} \frac{\frac{\partial y_{i,k}}{\partial v_{i,k}}}{\text{Neuron input}} \frac{\frac{\partial v_{i,k}}{\partial w_{i,j,k}}}{\text{Weights}}$$

- If activation function is logistic function σ then $\frac{\partial y_{i,k}}{\partial v_{i,k}} = \sigma(v_{i,k})(1 - \sigma(v_{i,k}))$
- Means you have a multiplier that maxes out at 0.25
- Max value is when $v_{i,k} = 0.5$
- $\sigma(0.5) = 0.25$
- As you back-propagate, you keep multiplying by this in DP... $.25 \times .25 \times .25 \dots$
gradient gets tiny
- Result is that output at time tick $t \dots$
 - ...has little interaction with input at tick $t - 100$ during back-propagation
 - Practically speaking: means back-propagation has limited-duration memory


- First proposed in 1997 by Hichreiter/Schimiduber
 - Uses a more complicated architecture in each recurrent unit
 - Now used in many state-of-the-art sequence-based ML
 - Translation, text processing, speech-to-text, many others...

No Vanishing Gradients?

- Key idea: have a state that passes from tick to tick
 - Explicitly decide which part of state to update
 - Which part of the state to forget
 - And which part of the state affects the output
- If “forget” is turned off
 - State just passes through un-modified
 - Derivative of identity function is 1
 - So gradient updates during back-propagation pass through un-modified
 - No vanishing gradients!

Basic LSTM Architecture

- An LSTM-based RNN is unrolled into a sequence of LSTM “modules”
 - Each module accepts input at a particular time tick
 - As well as state (“long term” memory) and last output (“short term” memory)
 - Uses those to output a view of its current state




Input (vector of floats each tick)


2

²Reprinted with permission. <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

A Note on Notation



- σ means a logistic activation layer at the top
- \tanh means a hyperbolic tangent activation layer at the top




Input (vector of floats each tick)

3

³Reprinted with permission. <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Now, Let's Go Over the Parts in Detail

- At the top is a state (long term memory) that passes through the unit
 - Possibly unimpeded
 - Possibly modified




^aReprinted with permission.

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Changing Long-Term Memory Contents

- There are two operations that can affect the state...
 - An item-by-item vector multiplication
 - Known as the “forget gate”
 - And an item-by-item vector addition
 - Known as the “input gate”




^aReprinted with permission.

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

At the Bottom...

- Memorized output and new input flow through the unit
 - They control the forget gate and the input gate
 - Might see new input and decide not to modify state
 - Or might decide to throw out old state and rebuild
- At the end...
 - Input and state of memory used to produce output




^aReprinted with permission.

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Now Let's Walk Through Input Processing

- First we run the “forget gate”
 - Last output and new input pass through a NN
 - With a logistic layer at top
 - Produces all values from 0 to 1
 - Item-by-item multiply with state




^aReprinted with permission.

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Forget Gate

- Controls what items in state are forgotten
 - Logistic produces a 0 for a dimension?
 - Then you will zero-out that dimension in the input




^aReprinted with permission.

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Next Is the “Input Gate”

- Input plus last output passes through two NNs
 - One with a tanh layer at the top... produces activations from -1 to +1
 - One with a logistic layer at the top... produces activations from 0 to 1
- Item-by-item multiply produces a vector of updates to state
 - This update is added into the long term memory
 - This is the “input gate”




^aReprinted with permission.

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Producing Output

- Now we have modified the state (long-term memory)
- Use the state to produce the output
 - First, push state through a tanh layer
 - Maps memory to values from -1 to +1




^aReprinted with permission.

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Output Gate

- And push the input and last output through a final NN
 - With a logistic activation at top
 - Produces values between 0 and 1
 - Item-by-item multiply with post-processed state
 - This decides what part of the state to output
 - Called the “output gate”



^aReprinted with permission.

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>


Output Typically Pushed Through a Final NN

- Output is simply a subset of the long-term memory
 - Like a brain dump
 - Simply reports encoded contents of the memory
 - Not directly useful to solve any task
 - Need to post-process (decode) it
- For example, if task is part-of-speech (POS) tagging
 - We might push output through a NN
 - Topped with a tanh layer
 - One neuron for each POS (noun, verb, adjective, etc.)
 - Push neuron outputs through a softmax
 - Output of softmax chooses final POS

Stacking These Modules

- Say we want some more learning capacity
 - Can “stack” LSTMs
- At each time tick
 - Have m stacked LSTM modules
 - Creates a stack with m layers
 - Each with its own set of parameters
 - LSTM in layer l has same internal NN params, no matter the time tick

Stacking LSTMs



Creates a Grid of LSTMs

- LSTM layer l at time tick t
 - Sends its state and output to LSTM module at layer l , time tick $t + 1$
 - Sends its output to LSTM module at layer $l + 1$, time tick t
- LSTM layer 1 at time tick t
 - Gets its input externally, from input stream
 - All other layers get their input from previous layer at tick t
- The output from the last/top layer at tick t ...
 - Is used to produce output at tick t

Uses for Stacked LSTM Modules

- Recognize more complex concepts/objects
- Learn hierarchies
- Predict activity from multiple signals

- ? How can we use what we learned today?
 - What LSTMs are and how they work
- ? What do we know now that we didn't know before?
 - Another way to handle sequential data in a neural network