

Tools & Models for Data Science

Deep Neural Networks (3): RNN

Chris Jermaine & Risa Myers

Rice University



Objectives

- Learn about another neural network architecture
- Learn about the types of problems which are well-suited to it

? What kinds of input have we considered so far (in the course)?

- How have we represented the text data?
 - Bag of words
- ? What are the limitations of this representation?

Issues with Data Representation

- How have we represented the text data?
 - Bag of words
- What are the limitations of this representation?
 - Lose the order/context
 - Fixed dictionary

Consider our Simple Feed-Forward Networks

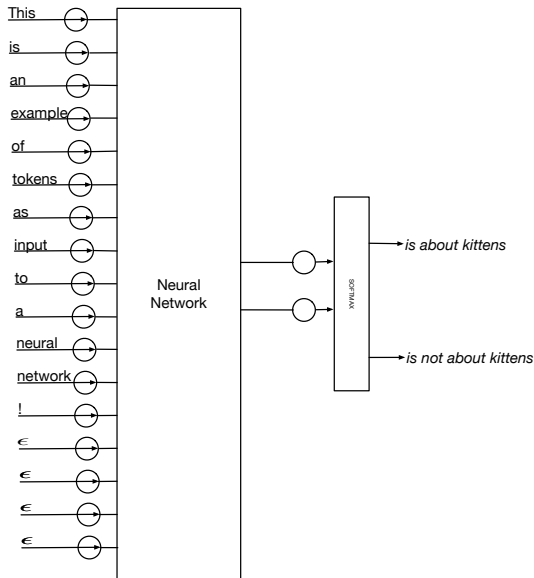
- They don't easily handle sequences
- What if I want to classify text docs?
 - And I don't want to do the bag-of-words thing
 - After all: bag-of-words loses word order
 - ? What are my options?

Consider our Simple Feed-Forward Networks

- They don't easily handle sequences
- What if I want to classify text docs?
 - And I don't want to do the bag-of-words thing
 - After all: bag-of-words loses word order
 - What are my options?
 - Sequence of tokens

Standard Idea

- Use FF network with enough input units (e.g. 20,000 tokens or 20,000 characters)
 - To handle any document in training
 - Pad unused tokens with a special character



- High model complexity
 - Max n input tokens
 - Size m first hidden layer
 - Means $n \times m$ weights to learn

$$10^4 \times 10^5 = 10^9 = 1\text{B weights}$$

? What if max tokens is 100K, average is 1000?

- High model complexity
 - Max n input tokens
 - Size m first hidden layer
 - Means $n \times m$ weights to learn
 - What if max tokens is 100K, average is 1000?
- ? What other issues are there?

- We are spending effort learning a network that is too big
- Likely not much training data for right-most inputs
- We can't handle bigger lengths (training or test)
- We have to use all the inputs

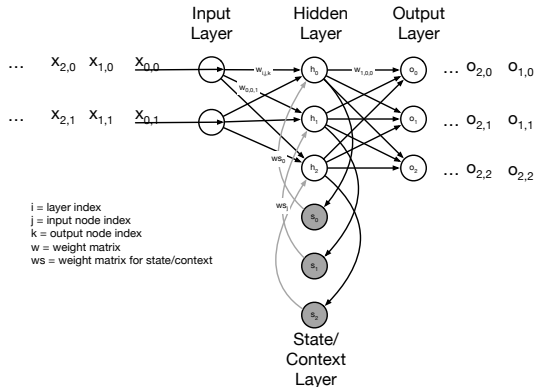
- Position i treated as different from position j
 - Not always the case!
 - If we see “kitten” at pos 34 or pos 1034, it is perhaps the same
 - We want to recognize pattern “kitten” regardless of position
 - Or consider swapping the order of 2 paragraphs on Wikipedia. No one would notice!
 - Whether or not it matters depends on the data / context

FFs on Fixed Length Sequences of Tokens

- Don't work well
- Alternatives?
 - Recurrent Neural Networks (RNNs)
- How do these help?
 - Add extra nodes that preserve context (order, state)
 - Incorporated via extra connections that cyclically link layers

What does an RNN look like?

- Classic RNN is Elman Network
- The output from the hidden layer is pushed to the output AND copied to the state layer
- Input from context neurons is fed in with the input data
- There are NO weights from the hidden layer neurons output to the states
- There ARE weights from the states to the neuron inputs
- There is a 1-1 correspondence between the saved state and each hidden node



- Has a set of context nodes (the "state layer") ("s" neurons in previous slide)
- They read value of the hidden layer
 - Non-trainable (e.g. values are not transformed)
 - Value simply remembered for one time tick
- To process t ticks of data:

```
init value of states in context/state layer to zeros
for  $i = 1$  to  $t$ 
  read input  $x_i$ 
  update hidden layer using  $x_i$  along with state layer
  if ( $i == t$ ) // for sequence-to-sequence, omit "if clause"
    hidden layer used to produce output
    hidden layer copied to state layer
end for
```

Data

- 1 Video data
- 2 Voice
- 3 Text
- 4 Other time sequence
 - Stock values
 - Temperatures

Input

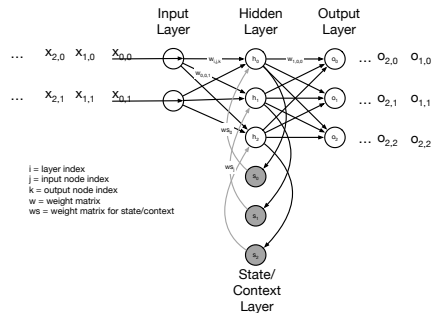
- 1 Sequence of images
- 2 Sequence of signal values
- 3 Sequence of characters/tokens
- 4 Sequence of [multidimensional] numeric values

Output

- Also a sequence
- Generated at each time tick
- Examples
 - 1 Video captions
 - 2 Translation
 - 3 Parts of speech
 - 4 Classification

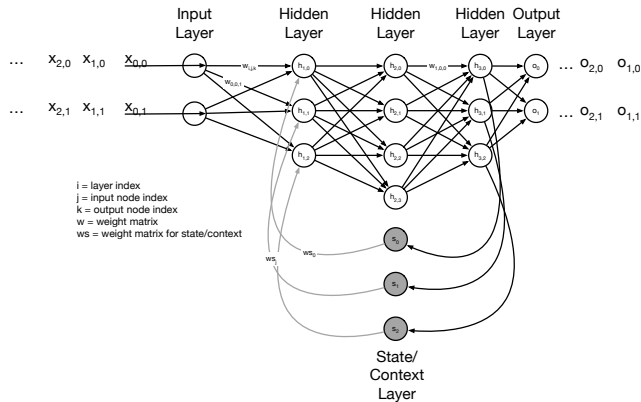
Elman Network in Practice

- Concatenate x_i and the state values
- Perform the matrix multiplication to get the inputs to the next layer
- At the top, use only the hidden layer to product the output



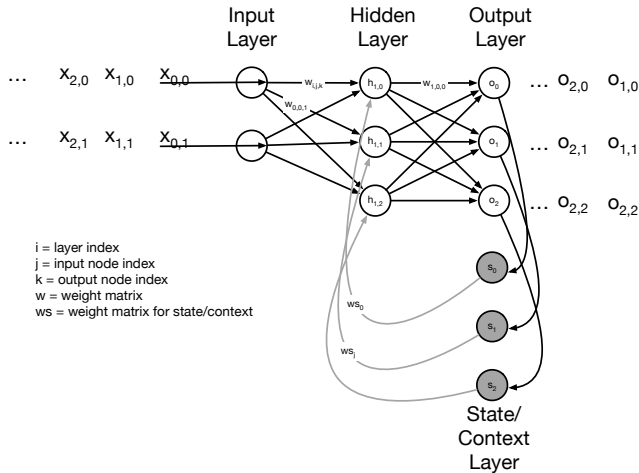
Elman Network

- We can have many hidden layers
- That is, a “deep net”
- In this case
 - Last hidden layer output copied to state
 - State used as input to first hidden layer...
 - In next time tick

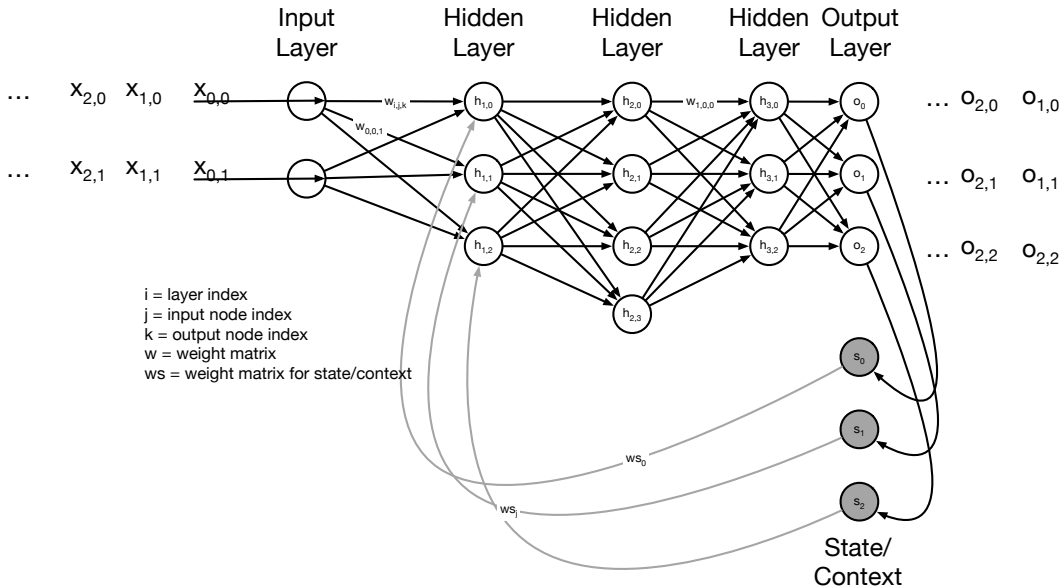


Jordan Network

- Similar, but copy output values, not hidden values
- Can be used for sequence-to-sequence
- Must be producing output at each tick

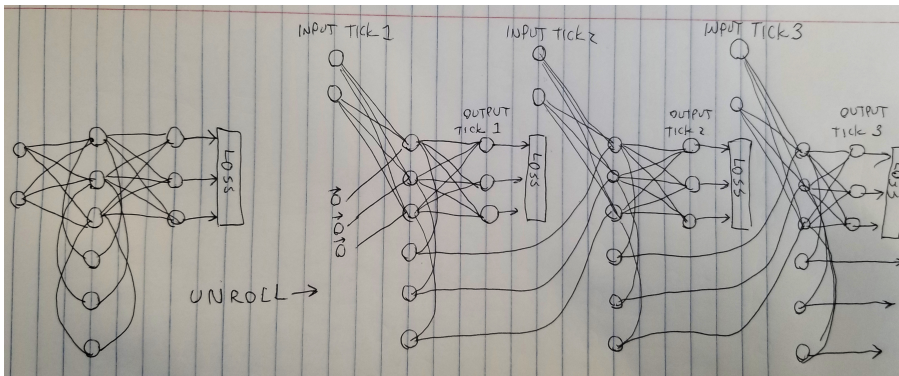


Deep Jordan Network



- Classic algorithm is back-propagation through time
- That is, view RNN as compact representation for a complex graph
- Unroll the complex graph
- And then use back-propagation on that
 - Key difference from classic back-propagation:
 - Weights are constrained to repeat

Example: Unrolling an Elman Network



- Example of unrolling a network for three time ticks
- Note distance backproped error from last time tick must travel
 - Goes through output neurons at time tick 3
 - Through hidden neurons time tick 3
 - Through hidden neurons time tick 2
 - Through hidden neurons time tick 1

Training Difficulty: Vanishing Gradient

- Errors fall off exponentially as backproped thru layers
 - Problem is that derivative of loss wrt activation function often $\ll 1$
 - Repeatedly multiplying causes backproped errors to tend to zero
 - Happens with deep, feed-forward nets, too
 - But unrolled RNNs are often especially deep

Training Difficulty: Vanishing Gradient

- Errors fall off exponentially as backproped thru layers
 - Problem is that derivative of loss wrt activation function often $\ll 1$
 - Repeatedly multiplying causes backproped errors to tend to zero
 - Happens with deep, feed-forward nets, too
 - But unrolled RNNs are often especially deep
- Means that in a deep net...
 - ...backprop does not affect weights much in first (leftmost) few layers

Training Difficulty: Vanishing Gradient

- Especially a problem if there is just one output at end of unrolled RNN
- Like in a pure classification task
 - Means that you will learn to classify
 - ...using only the last few time-tick's worth of data
 - Because early data can't interact with backproped error
 - Starting or ending data ends of being prioritized (depending on order data is fed in)
 - Mitigation
 - Batch normalization - normalize inputs to each layer

Problem With RNNs

- The “vanishing gradient” problem
- During back-propagation
 - Update magnitude drops exponentially with distance from output
 - Recall

$$\frac{\partial L}{\partial w_{i,j,k}} = \frac{\partial L}{\partial y_{i,k}} \frac{\partial y_{i,k}}{\partial v_{i,k}} \frac{\partial v_{i,k}}{\partial w_{i,j,k}}$$

- If activation function is logistic function σ then $\frac{\partial y_{i,k}}{\partial v_{i,k}} = \sigma(v_{i,k})(1 - \sigma(v_{i,k}))$
 - Means you have a multiplier that maxes out at 0.25
 - Max value is when $v_{i,k} = 0.5$
 - $\sigma(0.5) = 0.25$
 - As you back-propagate, you keep multiplying by this in DP... $.25 \times .25 \times .25 \dots$
gradient gets tiny
- Result is that output at time tick $t \dots$
 - ...has little interaction with input at tick $t - 100$ during back-propagation
 - Practically speaking: means back-propagation has limited-duration memory

- Very expensive to train
- Still expensive to use
- Alternatives?
 - Long Short Term Memory networks (LSTMs)

- Long Short Term Memory networks
 - Special RNN designed to deal with vanishing gradient problem
 - In LSTM, long term memory is not pushed through activation functions
 - So we don't have vanishing gradients
 - <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Tokens or characters?

- Depends on what you are trying to learn
- Tokens
 - Can stem words / tokens
 - Fewer parameters needed
 - Lower computational cost
- Characters
 - Smaller input dataset
 - Little to no preprocessing needed
 - Better for some languages (morphologically rich)
 - Grammar is reflected more in words than in position

- ? How can we use what we learned today?
 - What RNNs are and how they work
- ? What do we know now that we didn't know before?
 - How to handle sequential data in a neural network