Tools & Models for Data Science
Big Data Part Two: Beyond MapReduce

Chris Jermaine & Risa Myers

Rice University

Most Popular "Pure" MapReduce Software

- Is called Hadoop
  - Runs on JVM (like most Big Data software)
  - Includes MapReduce functionality
  - Plus the Hadoop distributed file system (HDFS)
- Hadoop popularity peaked around 2015...
- Has been declining since then
- Why? Were several issues...

- Data reread from DFS for each MR job
- Bad for iterative data processing
    - Example: most machine learning uses gradient descent
    - Need to make 100's or 1000's of passes over data
    - Re-evaluating gradient at various points

- Can only do Map
- Or MapReduce
- Everything else must be implemented in terms of those operations
- Unless you cheat
  - Ex: Mappers/Reducers talk to each other using sockets
  - But then why not just go with C/MPI?

Result: MapReduce Used Less and Less

- There are now an entire ecosystem of alternative softwares
  - Spark, Flink, etc.
- For use in both streaming and batch processing applications
- Generally oriented more towards in-memory computing
- Have far more expressive APIs
- We will focus on Spark

Apache Spark

- #1 Hadoop MapReduce killer
- What is Spark?
    - Platform for efficient distributed data analytics
- Runs on the JVM
- Written in Scala
    - But has bindings for Java, Scala, Python, R
    - Python nice for data analytics (NumPy, SciPy)... will focus there
- Doesn't do storage
    - Focus exclusively on compute
    - Commonly used with HDFS, S3, HBase, etc.

RDDs

- Basic abstraction: Resilient Distributed Data Set (RDD)
- Resilient - fault tolerance
- Distributed - across machines in a cluster
- Data Set - "working set of data"
- Read-only
- RDD is a data set buffered in RAM by Spark

- pySpark is interactive!
- RDDs
    - Are lazy: computations are done on demand and only as much as needed
    - Are ephemeral: computations are discarded when no longer used
    - Have lineage: the train of computations to generate an RDD is saved and may be replayed

Creating RDDs

■ To create and load an RDD:

```
data = [1, 2, 3, 4, 5]
myRDD = sc.parallelize (data)
```

■ Try it

Converting an RDD to a List

- collect
    - Be careful - this returns the ENTIRE RDD
- top(n)
    - Based on implicit ordering
- take(n)
    - First n elements

```
myRDD.collect()
myRDD.top(3)
myRDD.take(2)
```

- Try it

Create an RDD from a Range

- To create and load an RDD:

```
myRDD = sc.parallelize (range (20000))
```

- Try it

Other ways to create an RDD

- To create and load an RDD from a file:

```
myRDD = sc.textFile (someFileName) # sc is the Spark context
```

- Try it

Create an RDD from Another RDD

- Computations: Series of Xforms Over RDDs
- Example: word count

```
def countWords (fileName):
    lines = sc.textFile (fileName)
    tokens = lines.flatMap (lambda line: line.split(" "))
    instances = tokens.map (lambda word: (word, 1))
    aggCounts = instances.reduceByKey (lambda a, b: a + b)
    return aggCounts.top (200, key=lambda p: p[1])
```

- What transforms do we see here?
    - flatMap, map, reduceByKey, top
- Let's go through them

- But first, quick review of lambdas...
  - Fundamental to programming in Spark

What's a Lambda?

- Basically, a function that that we can pass like a variable
- Key ability: can "capture" its surroundings at creation
- Can also accept parameters

```python
def addTwelveToResult (myLambda):
    return myLambda (3) + 12

a = 23
aCoolLambda = lambda x : x + a
addTwelveToResult (aCoolLambda) # prints 38

a = 45
addTwelveToResult (aCoolLambda) # prints ???
```

- Try it

What's going on here?

- Anytime we see "myLambda" we replace it with the body of the lambda
- Kind of like a database VIEW

```
def addTwelveToResult (myLambda):
    return myLambda (3) + 12

a = 23
aCoolLambda = lambda x : x + a
addTwelveToResult (aCoolLambda)
```

- Results in

```
myLambda (3) + 12
(x + a) + 12
(3 + a) + 12
(3 + 23) + 12
```

- The parentheses matter!

- Lambdas can return many items
- Lambdas MUST return something

```
def sumThem (myLambda):
     tot = 0
     for a in myLambda ():
          tot = tot + a
     return tot

x = np.array([1, 2, 3, 4, 5])
iter = lambda : (j for j in x)
sumThem (iter) # prints 15
```

Spark Operation: flatMap ()

```
def countWords (fileName):
    lines = sc.textFile (fileName)
    tokens = lines.flatMap (lambda line: line.split("␣"))
```

- Processes every data item in the RDD
- Apply lambda to it
- Lambda argument will return zero or more results
- Can omit, combine or create elements
- Each result added into resulting RDD

Spark Operation: map ()

```
def countWords (fileName):
    lines = sc.textFile (fileName)
    tokens = lines.flatMap (lambda line: line.split("␣"))
    instances = tokens.map (lambda word: (word, 1))
```

- Process every data item in the RDD
- Apply lambda to it
- The lambda must return exactly one result
- The returned RDD has a new element with each element replaced by the lam

Spark Operation: reduceByKey ()

```
def countWords (fileName):
     lines = sc.textFile (fileName)
     tokens = lines.flatMap (lambda line: line.split(" "))
     instances = tokens.map (lambda word: (word, 1))
     aggCounts = instances.reduceByKey (lambda a, b: a + b)
```

- Data must be $(Key, Value)$ pairs
- Shuffle so that all (K, V) pairs with same K on same machine
- Organize into $(K, (V_1, V_2, ..., V_n))$ pairs
- Use the lambda to "reduce" the list to a single value
- This is similar to our aggregate functions in SQL!

Spark Operation: top ()

```
def countWords (fileName):
    lines = sc.textFile (fileName)
    tokens = lines.flatMap (lambda line: line.split("␣"))
    instances = tokens.map (lambda word: (word, 1))
    aggCounts = instances.reduceByKey (lambda a, b: a + b)
    return aggCounts.top (200, key=lambda p: p[1])
```

- Data must be $(Key, Value)$ pairs
- Takes two params... first is the number of list items to return
- Second (optional): lambda to use to obtain key for comparison
- Note: `top` collects an RDD, moving from cloud to local
- So result is **not** an RDD

An Important Note

- Spark uses lazy evaluation...
- If I run this code:

```
lines = sc.textFile (fileName)
tokens = lines.flatMap (lambda line: line.split("␣"))
instances = tokens.map (lambda word: (word, 1))
aggCounts = instances.reduceByKey (lambda a, b: a + b)
```

- Nothing happens! (Other than Spark remembers the ops)
  - Spark does not execute until an attempt made to collect an RDD
  - When we hit `top()`, then all of these are executed
- Why do this?
  - By waiting until last possible second, we can "pipeline"
  - Only operations that require a shuffle can't be pipelined

- groupByKey(), join(), reduce(), aggregate()

groupByKey ()

- Data must be $(Key, Value)$ pairs
- Shuffle so that all (K, V) pairs with same K onto the same machine
- Organize into $(K, \langle V_1, V_2, ..., V_n \rangle)$ pairs
- Store each list as a `ResultIterable` for future processing
- Like `reduceByKey()` but without the reduce

join ()

- Given two data sets *rddOne*, *rddTwo* of (*Key*, *Value*) pairs
- We join them using:

```
rddOne.join (rddTwo)
```

- Returns $(K, (V_1, V_2))$ pairs
- Constructed from all (K_1, V_1) from *rddOne*, (K_2, V_2) from *rddTwo*, where $K_1 = K_2$

- Example:
  - *rddOne* is $\{(red, 9), (blue, 7), (red, 12), (green, 4)\}$
  - *rddTwo* is $\{(blue, up), (green, down), (green, behind)\}$
  - Result of join is $\{(blue, (7, up)), (green, (4, down)), (green, (4, behind))\}$
- Can blow up RDD size if join is many-to-many
- Requires expensive shuffle!
- ? How do we write the join?

reduce()

- Unlike past operations, this is not a transform from RDD to RDD
  - This is like `top()`
  - It moves the result back to Python
- Repeatedly applies a lambda to each item in the RDD to get single result

```
>>> myData = sc.parallelize (range(20000))
>>> myData.reduce (lambda a, b: a + b)
199990000
```

aggregate()

- With `reduce()` you aggregate directly, can be restrictive...
- Example: RDD is $\{(red, 9), (blue, 7), (red, 12), (green, 4)\}$
  - Want: **dictionary where value is sum for each unique color**
  - Cannot use `reduce()`
  - It only "sums" up the items in the input RDD directly
  - Two inputs and output must be the same type
  - How do I get the desired output dictionary by defining $+$ in $(((red, 9) + (blue, 7)) + (red, 12)) + (green, 4)$?

aggregate() (Continued)

- Data must be $(Key, Value)$ pairs
- Organize into $(K, \langle V_1, V_2, ..., V_n \rangle)$ pairs
- Then aggregate the list, like `reduce()`
- `aggregate()` takes three arguments
    - The "zero" to init the aggregation
    - Lambda that takes X_1, X_2 and aggregates them, where X_1 already aggregated, X_2 not
    - Lambda that takes X_1, X_2 and aggregates them, where both aggregated

aggregate() (Example)

```
1  def add (dict, tuple):
2    result = {}
3    for key in dict:
4      result[key] = dict[key]
5    if (tuple[0] in result):
6      result[tuple[0]] += tuple[1]
7    else:
8      result[tuple[0]] = tuple[1]
9    return result
```

Goal: Add tuple to dict and
return a new dictionary

1. Define add that takes a dictionary and a tuple
2. initialize the result dictionary
3. For each item in the dictionary we are given
4. Create an entry in our result dictionary with dict's value
5. If we've seen the tuple key before
6. Accumulate the value
7. If this is the first time we see the tuple key
8. Create an entry for it and initialize with the value
9. Return our result dictionary

aggregate() (Example)

```
1   def combine (dict1, dict2):
2     result = {}
3     for key in dict1:
4       result[key] = dict1[key]
5     for key in dict2:
6       if (key in result):
7         result[key] += dict2[key]
8       else:
9         result[key] = dict2[key]
10    return result
```

Goal:

Return a new dictionary that contains all the keys in dict1 and dict2 with the total counts

1. Define combine that takes 2 dictionaries
2. initialize the result dictionary
3. For each item in dict1 that we are given
4.   Create an entry in our result dictionary with dict1's value
5. For each key, value in dict2 that we are given
6.   If we've seen the key before
7.   Accumulate the value
8.   If this is the first time we see the tuple key
9.   Create an entry for it and initialize with the value
10. Return our result dictionary

aggregate() (Example)

```
>>> myRdd = sc.parallelize ([('red', 9), ('blue', 7),
... ('red', 12), ('green', 4)])
>>> myRdd.aggregate ({}, lambda x, y: add (x, y),
... lambda x, y: combine (x, y))

{'blue': 7, 'green': 4, 'red': 21}
```

Closing Thoughts

- When is Spark/MapReduce a better option than HPC?
  - When your pipeline is heavily data-oriented
  - Or when your compute is (relatively) loosely coupled
- Key benefits compared to HPC
  - Built in fault tolerance
  - Better support for BIG data
  - Much higher programmer productivity
- Will continue to take market share from HPC
  - You see academic papers with both MPI, Spark implementations
  - But not everything can move to Spark

Wrap up

```
https:
//spark.apache.org/docs/latest/rdd-programming-guide.html
```

```
https://spark.apache.org/docs/latest/api/python/
```

? How can we use what we learned today?

? What do we know now that we didn't know before?